

# Suryansh Singh Rawat

+1 (206) 496 5588 ✉ [suryansh@uw.edu](mailto:suryansh@uw.edu) [in suryansh-singh-rawat](https://www.linkedin.com/in/suryansh-singh-rawat) [xsuryanshx](https://x.com/suryanshx) [snareyansh](https://www.snoow.com/u/sureyansh) [suryansh.space](https://suryansh.space)

## OBJECTIVE

AI/ML Engineer with 4 years of experience building enterprise-scale RAG, agentic LLM systems and cybersecurity NLP research, focused on shipping production-grade AI systems at scale. Currently Seeking Summer 2026 MLE / AI Research Intern roles.

## EDUCATION

### University of Washington, Seattle

Sept 2025 – Mar 2027

Master of Science (M.S), Data Science | **GPA: 3.81/4.0**

Seattle, WA

- Relevant Courses: Natural Language Processing, Machine Learning, Probability and Statistics, Applied Statistics and Experimental Design, Software Design for Data Science, Scalable Data Systems and Algorithms

### Birla Institute of Technology and Science, Pilani

Aug 2018 – July 2022

Bachelor of Engineering (B.E), Electronics and Instrumentation

Goa, India

## TECHNICAL SKILLS

**Languages:** Python, C++, SQL, R, Java, TypeScript, Bash, MATLAB

**Developer Tools:** Git, AWS, Microsoft Azure, Docker, S3, DynamoDB, PostgreSQL, MongoDB, Redis, Supabase

**Frameworks:** PyTorch, TensorFlow, Keras, SpaCy, HuggingFace, LangChain, LangGraph, Streamlit, FastAPI, Selenium

## WORK EXPERIENCE

### Ascentt

Mar 2024 – Sept 2025

Principal AI Engineer

Remote, US

- Scaled ToyotaGPT (ChatGPT equivalent for enterprises) to **74,000+** active users, architecting **REST APIs**, **DynamoDB** for user state management, and an **Agentic RAG** pipeline using **LangGraph** managing **1TB+ of unstructured data** across **14+** document types.
- Optimized RAG pipelines across **4 enterprise** products, implementing lossless preprocessing, chunking, and indexing in **Amazon OpenSearch**, and benchmarking with **RAGAS**, achieving **+25% Precision@10** and **+12% Recall@10** over vanilla RAG baseline.
- Owned secure file ingestion for enterprise AI systems, designing **encrypted upload** and preprocessing pipelines using **AES-GCM with AWS KMS**, enabling compliant handling of sensitive documents across multiple file types.
- Led backend repository maintenance** and code reviews, ensuring high-quality, production-ready code across enterprise systems.

### Deloitte USI

July 2022 – Mar 2024

Data Scientist | AI Center of Excellence Team

Bangalore, India

- Developed **Autoencoder anomaly detection** models with graph-based preprocessing in **Neo4J** to detect **zero-day threat attacks** from real-time AWS network flow data.
- Researched an NLP pipeline for regulatory document consolidation, leveraging **metric learning** trained **SecBERT** embeddings and **HDBSCAN** to merge semantically overlapping compliance requirements, achieving **~83% reduction** in manual SME efforts.
- Constructed a **SecBERT text summarization model** for generating cyber threat reports, reducing analyst review time by **~40%**, integrating **NER tagging and masking** with **SpaCy** annotation, deploying end-to-end pipeline using **Flask and Docker**.

### Scienaptic AI

July 2021 – Dec 2021

Data Scientist Intern

Bangalore, India

- Achieved an **11.5% increase** in application approvals by building credit underwriting models using **XGBoost and Logistic Regression** to forecast the probability of credit accounts that may default.
- Built a unified credit underwriting model using **100M+ raw credit records** to deliver baseline metrics across diverse risk segments.

## SELECTED PROJECTS

### OpenProbe | Python, Langchain, LangGraph

May 2025 – June 2025

- Built OpenProbe, a deep research **multi agent orchestration** using **LangGraph** state orchestration (websearch, code, reasoning).
- Achieved **67.1% accuracy** on the FRAMES benchmark, outperforming GPT-4o-Search by **1.5%**, using DeepSeek-R1/Qwen3-32B.

### Ad-Astra: xAI Ad Intelligence Engine | Python, Grok, XAPI, Supabase

Dec 2025 – Dec 2025

- Built a personalized ad generation and placement engine for X (Twitter), using **Grok 4.1 Reasoning & Vision models** for persona-aware ad copy generation and **scroll telemetry** to optimize in-feed placement; improving engagement and ad relevance.

### PersonaAI | Python, Microsoft Phi-4, Supabase

Oct 2025 – Oct 2025

- Built AI platform **simulating A/B tests** on **1M+ NVIDIA Nemotron Personas** dataset, to predict digital campaign performance.
- Created natural-language audience targeting with variant-level diagnostics; **3rd place at Lovable Hackathon** against **40+** teams.

### KeyCognition | Python, PyTorch, React

Nov 2025 – Nov 2025

- Built a behavioural analysis platform processing **136M+ keystroke events** to predict user cognitive load (low/med/high).
- Leveraged temporal typing features (e.g., spelling mistakes, backspaces) with **LSTM and K-Means models** for user behavior classification, winning **3rd place at the UW Databricks Hackathon**.

### Detecting GAN Generated DeepFake Images | Python, Tensorflow, Deep Learning

Jan 2021 – June 2021

- Designed lightweight CNN architecture to detect **GAN-generated** DeepFakes, achieved **97.77% accuracy** on StyleGAN dataset with SOTA-level precision and recall.