Suryansh Singh Rawat

+1 (206) 496 5588 suryansh@uw.edu suryansh-singh-rawat xsuryanshx snareyansh suryansh.space

OBJECTIVE

AI/ML Engineer with 4 years of experience building enterprise-scale RAG, agentic LLM systems and cybersecurity NLP research, focused on shipping production-grade AI systems at scale. Currently Seeking Summer 2026 MLE / AI Research Intern roles.

EDUCATION

University of Washington, Seattle **Sept 2025 – Mar 2027**
Master of Science (M.S), Data Science | ***GPA: 3.81/4.0*** *Seattle, WA*

- Relevant Courses: Natural Language Processing, Machine Learning, Probability and Statistics, Applied Statistics and Experimental Design, Software Design for Data Science, Scalable Data Systems and Algorithms

Birla Institute of Technology and Science, Pilani **Aug 2018 – July 2022**
Bachelor of Engineering (B.E), Electronics and Instrumentation *Goa, India*

WORK EXPERIENCE

Ascentt **Mar 2024 – Sept 2025**
Principal AI Engineer *Remote, US*

- Scaled ToyotaGPT (ChatGPT equivalent for enterprises) to **74,000+** active users, architecting **REST APIs**, **DynamoDB** for user state management, and an **Agentic RAG** pipeline using **LangGraph** managing **1TB+ of unstructured data** across **14+** document types.
- Optimized RAG pipelines across **4 enterprise** products, implementing lossless preprocessing, chunking, and indexing in **Amazon OpenSearch**, and benchmarking with **RAGAS**, achieving **+25% Precision@10 and +12% Recall@10** over vanilla RAG baseline.
- Owned secure file ingestion for enterprise AI systems, designing **encrypted upload** and preprocessing pipelines using **AES-GCM with AWS KMS**, enabling compliant handling of sensitive documents across multiple file types.
- **Led backend repository maintenance** and code reviews, ensuring high-quality, production-ready code across enterprise systems.

Deloitte USI **July 2022 – Mar 2024**
Data Scientist | AI Center of Excellence Team *Bangalore, India*

- Developed Deep Learning **Autoencoder architectures** with **Graph Network** based preprocessing in **Neo4J** to identify anomalies for detecting **Zero-Day Cyber Threats** using realtime AWS cloud network flow data.
- Researched an NLP pipeline for regulatory document consolidation, leveraging **metric learning** trained **SecBERT** embeddings and **HDBSCAN** to merge semantically overlapping compliance requirements, achieving ∼**83% reduction** in manual SME efforts.
- Constructed a **SecBERT text summarization model** for generating cyber threat reports, reducing analyst review time by ∼**40%**, integrating **NER tagging and masking** with **SpaCy** annotation, deploying end-to-end pipeline using **Flask and Docker**.

Scienaptic AI **July 2021 – Dec 2021**
Data Scientist Intern *Bangalore, India*

- Achieved an **11.5% increase** in application approvals by building Credit Underwriting ML Models (XGBoost and Logistic Regression) to forecast the probability of credit accounts that may default.
- Built a unified credit underwriting framework using **100MM+ raw credit records** to deliver baseline metrics across risk segments.

KEY TECHNICAL PROJECTS

OpenProbe | *Python, Langchain, LangGraph* **May 2025 – June 2025**

- Built OpenProbe, an open-source deep research agent using LangGraph state orchestration (websearch, code execution, reasoning).
- Outperforms GPT-4o-Search on multi-hop FRAMES benchmark achieving 67.1% accuracy with DeepSeek-R1/Qwen3-32B.

Ad-Astra: xAI Ad Intelligence Engine | *Python, Grok, XAPI, Supabase* **Dec 2025 – Dec 2025**

- Developed a hyper-personalized advertising engine for X (Twitter) that generates tailored ad copies based on user persona and utilizes scroll telemetry for intelligent in-feed placement, improving user engagement and ad relevance across simulated platform users.

KeyCognition | *Python, PyTorch, React* **Nov 2025 – Nov 2025**

- Built a behavioural analysis platform processing 136M+ keystroke events to predict user cognitive load (low/med/high).
- Leveraged temporal typing features (eg. spelling mistakes, using backspace) and LSTM + K-Means models for real-time user behavior classification.

Detecting GAN Generated DeepFake Images | *Python, Tensorflow, Deep Learning* **Jan 2021 – June 2021**

- Designed lightweight CNN architecture to detect GAN-generated DeepFakes, achieved 97.77% accuracy on StyleGAN dataset with SOTA-level precision and recall.

TECHNICAL SKILLS

Languages: Python, SQL, C++, Java, TypeScript, Bash, MATLAB, R
Developer Tools: Git, AWS, Microsoft Azure, Docker, S3, DynamoDB, PostgreSQL, MongoDB, Redis, Supabase
Frameworks: PyTorch, TensorFlow, Keras, SpaCy, HuggingFace, CUDA, LangChain, LangGraph, Streamlit, FastAPI, Selenium